

Advanced Control Architectures for Autonomous Systems

Lecture 8: Motion Planning Algorithms (RRT*, RRT, DWA, path generation)

Autonomous Mobile Robots

Fall 2025

Project Expectations

Three Essential Components

① System Modeling

- Kinematics, dynamics, mathematical formulation
- Physical working principles and implementation

② Control Methodology

- Justification of chosen approach
- Comprehensive testing and validation

③ Motion Planning

- Integration of planning algorithms
- Explanation of underlying principles

Deterministic vs Probabilistic Control

Fundamental Philosophy

Deterministic methods rely on known physics, not probability distributions

Deterministic Control (NMPC)

- System model:
 $\dot{x} = f(x, u)$
- Physics is reliable and known
- Future states are predictable
- Suitable for well-modeled systems

Probabilistic Methods (RL)

- System: $P(s'|s, a)$ (MDPs)
- Environment is uncertain
- Outcomes are stochastic
- Optimizes expected reward

Selection criterion: Confidence in physical model accuracy

Case Study: Frozen Lake Problem

Stochastic Environment Scenario

Scenario: Robot navigation on slippery ice

Command: "Move Forward"

Actual outcomes:

- 33% probability: Move Forward
- 33% probability: Slide Left
- 33% probability: Slide Right

Limitation: NMPC assumes deterministic dynamics

Solution: Reinforcement learning handles probability distributions

NMPC Optimization Formulation

Optimal Control Problem

At each time step t , solve over prediction horizon T_p :

$$\min_u J = \int_t^{t+T_p} [\|x(\tau) - x_{ref}(\tau)\|_Q^2 + \|u(\tau)\|_R^2] d\tau + \|x(t + T_p) - x_{goal}\|_P^2$$


Subject to constraints:

- System dynamics: $\dot{x} = f(x, u)$
- State constraints: $x \in \mathcal{X}_{safe}$
- Input constraints: $u \in \mathcal{U}_{feasible}$

Prediction Horizon

T_p determines how far into the future the controller predicts

Receding Horizon Principle



Control Execution

- ① Measure current state $x(t)$
- ② Solve optimization over horizon $T_p \rightarrow$ obtain $u^*(t)$
- ③ Apply only first control $u^*(t)$
- ④ Shift window and repeat

Cost Function Components

Tracking Error	Control Effort	Terminal Cost
$\ x - x_{ref}\ _Q^2$	$\ u\ _R^2$	$\ x(T_p) - x_{goal}\ _P^2$
<ul style="list-style-type: none">• Penalizes path deviation• Large Q: tight tracking• Small Q: loose tracking	<ul style="list-style-type: none">• Penalizes large inputs• Large R: smooth control• Small R: aggressive	<ul style="list-style-type: none">• Ensures final state quality• Improves stability• Guides to goal region

Engineering Insight

Weights Q , R , P are tuning parameters - balance tracking vs control effort

Real-World Application: Autonomous Racing

Abu Dhabi Autonomous Racing League

- **Winner:** Technical University of Munich (NMPC implementation)
- **Performance:** 250+ km/h, zero crashes during race
- **Critical:** Tire friction modeling $F_{lateral} = \mu(v, \alpha) \cdot F_N$

Key Insight

"At 250 km/h, PID is reactive - by the time you see an error, you've already crashed. You NEED prediction."

Gradient Descent Solver

Continuous-Time Update Law


$$\dot{u} = -\alpha \nabla J(u) = -\alpha \frac{\partial J}{\partial u}$$

- \dot{u} represents **algorithmic time derivative**
- Solver "flows" down cost surface toward minimum
- Not physical system time - this is solver iteration

Implementation

Gradient descent operates in solver time, not physical time

Visualizing Gradient Descent



Process

- Start with initial guess u_0
- Compute gradient direction
- Step opposite to gradient: $u_{k+1} = u_k - \alpha \nabla J$
- Converge to local minimum

Lyapunov Stability Analysis

Proof of Convergence

Step 1 - Lyapunov Candidate: $V(u) = J(u)$

Step 2 - Time Derivative:

$$\begin{aligned}\dot{V} &= \frac{dJ}{dt} = \frac{\partial J}{\partial u} \cdot \frac{du}{dt} \\ &= \frac{\partial J}{\partial u} \cdot \left(-\alpha \frac{\partial J}{\partial u} \right) \\ &= -\alpha \left\| \frac{\partial J}{\partial u} \right\|^2\end{aligned}$$

Step 3 - Stability: $\dot{V} = -\alpha \|\nabla J\|^2 \leq 0$

Guarantee

Cost decreases monotonically \rightarrow Solver converges stably

Numerical Gradient Estimation

Finite Difference Approximation

$$\frac{\partial J}{\partial u} \approx \frac{J(u + \delta u) - J(u)}{\delta u}$$

Implementation Steps

- ① Simulate with u , compute $J(u)$
- ② Perturb: $u' = u + \delta u$
- ③ Simulate with u' , compute $J(u')$
- ④ Gradient: $\nabla J \approx \frac{J(u') - J(u)}{\delta u}$
- ⑤ Update: $u_{new} = u - \alpha \nabla J$

Critical Parameter

δu must be small but not too small (typical: 10^{-6})

Critical Simulation Rule

10× Bandwidth Rule

$$\Delta t_{sim} \leq \frac{1}{10} \times \tau_{system}$$


Practical Example

- Motor time constant: $\tau = 0.1$ seconds
- Required: $\Delta t \leq 0.01$ seconds
- Prediction horizon $T_p = 2s \rightarrow 200$ simulation steps

Consequences of Violation

- Numerical integration errors
- Simulation divergence
- Controller instability
- "Robot predicts it's in another country!"

Navigation Hierarchy



Dynamic Window Approach (DWA)

Core Principle

Search in **velocity space** (v, ω) instead of position space (x, y)

Dynamic Window

$$V = [V_t - a_{min} \cdot dt, V_t + a_{max} \cdot dt]$$


$$\Omega = [\omega_t - \alpha_{min} \cdot dt, \omega_t + \alpha_{max} \cdot dt]$$

- Based on acceleration limits
- Defines reachable velocities
- Ensures feasibility

Advantage

- Pre-computed trajectories
- Fast minimization
- Real-time performance

Visualizing DWA Search Space



Multi-Objective Optimization

$$G(v, \omega) = \alpha \cdot \text{Heading}(v, \omega) + \beta \cdot \text{Clearance}(v, \omega) + \gamma \cdot \text{Velocity}(v, \omega)$$

- **Heading:** Alignment with goal direction
- **Clearance:** Distance to nearest obstacle
- **Velocity:** Progress toward goal

Typical Configuration

10 linear velocities \times 10 angular velocities = 100 trajectories to evaluate

Rapidly-exploring Random Trees (RRT)


Algorithm

- ① **Sample:** Random state x_{rand} in free space
- ② **Nearest:** Find closest node x_{near} in tree
- ③ **Steer:** Extend from x_{near} toward x_{rand} by Δq
- ④ **Check:** If collision-free, add x_{new} to tree
- ⑤ **Repeat:** Until goal reached or maximum iterations

Key Feature

Rapidly explores high-dimensional spaces by random sampling

Visualizing RRT Expansion



Exploration Strategy

Tree grows toward randomly sampled points, avoiding obstacles

Computational Efficiency: KD-Trees

Complexity Challenge

- Naive nearest neighbor: $O(N)$ for N nodes
- With KD-tree: $O(\log N)$ complexity
- Example: $N = 10,000 \rightarrow 14$ comparisons vs 10,000

KD-Tree Structure

- Binary space partitioning
- Alternating axis splits
- Essential for real-time RRT

Key Point

"KD-trees are essential! Without them, RRT is too slow for real-time."

RRT vs RRT*: Optimality

Standard RRT

- Finds *any* feasible path
- Fast exploration
- No optimality guarantees
- Good for: Quick solutions


RRT* (Optimal)

- Asymptotically optimal
- "Rewiring" improves paths
- Converges to optimum
- Good for: Quality paths

RRT* Rewiring

- Check if X_{new} provides better paths to neighbors
- Rewire tree to reduce costs
- Continuous improvement over time

Complete System Architecture



Integrated Operation

- Global planning for strategic navigation
- Local adaptation for dynamic obstacles
- Physics-based control for precise execution
- Continuous sensor feedback for real-time updates

Key Technical Insights

① Deterministic Control Foundation

- NMPC relies on accurate physical models
- Gradient descent with Lyapunov stability guarantees
- Suitable for well-characterized systems

② Critical Implementation Rules

- $10 \times$ bandwidth rule: $\Delta t \leq \tau_{\text{system}}/10$
- Proper finite difference gradient estimation
- Careful weight selection in cost functions

③ Hierarchical Planning

- Global: RRT/RRT* for strategic paths
- Local: DWA for dynamic obstacle avoidance
- Control: NMPC for physics-based execution

Practical Implementation Guidelines

Development Strategy

1. Start Simple

- Basic controller implementation first
- Incremental complexity addition
- Early testing and validation

2. Method Selection

- Match complexity to application needs
- Classical control often sufficient
- Advanced methods for demanding applications

3. Validation Process

- Comprehensive simulation before hardware
- Careful time step verification
- Iterative weight tuning

Recommended Resources

Reference Materials

- **Planning Algorithms:** LaValle, "Planning Algorithms"
- **RRT Foundation:** LaValle (1998)
- **Optimal RRT:** Karaman & Frazzoli, "RRT*"
- **Digital Control:** Various sampling theory texts

Questions?